

# Developing a fetal classifier.

---

A data science project by  
Kenneth Imade

Photo by Volodymyr Hryshchenko on Unsplash





# Objective

- To develop a machine learning model capable of classifying fetus health into three categories; normal, suspect, or pathological
- Evaluate the performance of various machine learning classifiers and choose the best one for classifying fetus health

# Overview

## Context

Reduction of child mortality is reflected in several of the United Nations' Sustainable Development Goals and is a key indicator of human progress. The UN expects that by 2030, countries end preventable deaths of newborns and children under 5 years of age, with all countries aiming to reduce under-5 mortality to at least as low as 25 per 1,000 live births.


Parallel to notion of child mortality is of course maternal mortality, which accounts for 295 000 deaths during and following pregnancy and childbirth (as of 2017). The vast majority of these deaths (94%) occurred in low-resource settings, and most could have been prevented.

In light of what was mentioned above, Cardiotocograms (CTGs) are a simple and cost accessible option to assess fetal health, allowing healthcare professionals to take action in order to prevent child and maternal mortality. The equipment itself works by sending ultrasound pulses and reading its response, thus shedding light on fetal heart rate (FHR), fetal movements, uterine contractions and more.

**Data Source:** [Fetal Health Classification | Kaggle](#)

**Acknowledgement:** Ayres de Campos et al. (2000) SisPorto 2.0 A Program for Automated Analysis of Cardiotocograms. J Matern Fetal Med 5:311-318 (link)

# Exploratory Data Analysis



# Dataset Statistics

<b>N = 2,126</b>	<b>Mean</b>	<b>Std</b>	<b>Min</b>	<b>25%</b>	<b>50%</b>	<b>75%</b>	<b>Max</b>
Baseline value	133.303857	9.840844	106.0	126.000	133.000	140.000	160.000
Accelerations	0.003178	0.003866	0.0	0.000	0.002	0.006	0.019
Fetal movement	0.009481	0.046666	0.0	0.000	0.000	0.003	0.481
Uterine contractions	0.004366	0.002946	0.0	0.002	0.004	0.007	0.015
Light decelerations	0.001889	0.002960	0.0	0.000	0.000	0.003	0.015
Severe decelerations	0.000003	0.000057	0.0	0.000	0.000	0.000	0.001
Prolongued decelerations	0.000159	0.000590	0.0	0.000	0.000	0.000	0.005
Abnormal short term variability	46.990122	17.192814	12.0	32.000	49.000	61.000	87.000
Mean value of short term variability	1.332785	0.883241	0.2	0.700	1.200	1.700	7.000
Percentage of time with abnormal long term variability	9.846660	18.396880	0.0	0.000	0.000	11.000	91.000
Mean value of long term variability	8.187629	5.628247	0.0	4.600	7.400	10.800	50.700


# Dataset Statistics Continued

<b>N = 2,126</b>	<b>Mean</b>	<b>Std</b>	<b>Min</b>	<b>25%</b>	<b>50%</b>	<b>75%</b>	<b>Max</b>
Histogram width	70.445908	38.955693	3.0	37.000	67.500	100.000	180.000
Histogram min	93.579492	29.560212	50.0	67.000	93.000	120.000	159.000
Histogram max	164.025400	17.944183	122.0	152.000	162.000	174.000	238.000
Histogram number of peaks	4.068203	2.949386	0.0	2.000	3.000	6.000	18.000
Histogram number of zeroes	0.323612	0.706059	0.0	0.000	0.000	0.000	10.000
Histogram mode	137.452023	16.381289	60.0	129.000	139.000	148.000	187.000
Histogram mean	134.610536	15.593596	73.0	125.000	136.000	145.000	182.000
Histogram median	138.090310	14.466589	77.0	129.000	139.000	148.000	186.000
Histogram variance	18.808090	28.977636	0.0	2.000	7.000	24.000	269.000
Histogram tendency	0.320320	0.610829	-1.0	0.000	0.000	1.000	1.000
<b>Fetal health (Target Variable)</b>	<b>1.304327</b>	<b>0.614377</b>	<b>1.0</b>	<b>1.000</b>	<b>1.000</b>	<b>1.000</b>	<b>3.000</b>

The background features a target with concentric red and grey rings on the right side. On the left, several grey arrows point towards the target, and a single red arrow points directly at the center bullseye. The text is overlaid on this graphic.

What is the distribution  
of the target variable?

## Most Fetuses are classified as normal






The background features a dark blue-grey field. On the left, several grey arrows of varying lengths point horizontally towards the right. A single, prominent red arrow also points right, positioned behind the text. On the right side, there is a target graphic consisting of concentric circles in shades of red and grey. The text is centered horizontally and overlaid on the arrows.

What is the relationship  
between the target variable  
and other variables?


Strong correlation exists between fetal health and other factors






How does Fetal health change  
with respect to these  
variables?


**Tendency for both normal and suspect fetuses movement to slightly increase in tandem with their acceleration but relationship flatlines with pathological fetuses**




**Increased short term fetal heart rate suggests a decrease in fetal movement across all fetal health categories**



**Strong relationship between long term variability and fetal movement in pathological fetuses but relationship is weaker in suspect fetuses and insignificant in normal fetuses**



**Fetal movement increases alongside prolonged decelerations in both normal and suspect fetuses but does not for pathological fetuses**









# Data preprocessing and model building



**Features in dataset have varied ranges**



**Feature ranges standardised for model building**



# Choosing the best model

```
lr_model = Pipeline([('lr_classifier', LogisticRegression(random_state=42))])
dt_model = Pipeline([('dt_classifier', DecisionTreeClassifier(random_state=42))])
rf_model = Pipeline([('rf_classifier', RandomForestClassifier(random_state=42))])
knn_model = Pipeline([('knn_classifier', KNeighborsClassifier())])
svc_model = Pipeline([('svc_classifier', SVC())])

pipelines = [lr_model, dt_model, rf_model, knn_model, svc_model]
pipe_dict = {0: 'Logistic Regression', 1: 'Decision Tree', 2: 'Random Forest', 3: 'KNN', 4: 'SVC'}

for pipe in pipelines:
    pipe.fit(X_train, y_train)

#cross validation on accuracy
cv_results_accuracy = []
for i, model in enumerate(pipelines):
    cv_score = cross_val_score(model, X_train, y_train, cv=10)
    cv_results_accuracy.append(cv_score)
    print("%s: %f " % (pipe_dict[i], cv_score.mean()))
```

# Classification models accuracy result



# Models Accuracy Result


Logistic Regression: 0.897170

Decision Tree: 0.916683

Random Forest: 0.938187

KNN: 0.894504

# Hyper-parameter tuning




```
parameters = {
    'n_estimators': [100, 150, 200, 500, 700, 900],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [4, 6, 8, 12, 14, 16],
    'criterion': ['gini', 'entropy'],
}

rfc_cv = GridSearchCV(estimator=RandomForestClassifier(), param_grid=parameters, cv= 5)
rfc_cv.fit(X_train, y_train)

rfc_cv.best_params_
```

# Results

---




# Classification Metrics



	Precision	Recall	F1-Score	Support
Normal (1.0)	0.96	0.98	0.97	496
Suspect (2.0)	0.92	0.78	0.84	101
Pathological (3.0)	0.88	0.93	0.90	41
Accuracy			0.95	638
Macro Avg	0.92	0.90	0.91	638
Weighted Avg	0.95	0.95	0.95	638

**Confusion Matrix**





# Conclusion

The best model was found to be the **Random Forest Classifier**. The model is most accurate with the 'Normal' class, with a high number of correct predictions and a relatively small number of false predictions. The model performs reasonably well on the 'Pathological' class. The 'Suspect' class has the most misclassifications relative to its true count, with a notable number of false negatives (predicted as 'Normal').